

WHEN DOES LEAVE-ONE-OUT SUFFICE? AUDITING MINIMAL-CAUSE ATTRIBUTION IN RAG

Qing Lin¹, Long Zhang², Xudong Zhang^{3,*}

¹Dalian University of Technology

²Inspur Cloud Information Technology Co., Ltd., Jinan, Shandong, China

³Institute of Software, Chinese Academy of Sciences, Beijing, China

linq@mail.dlut.edu.cn; zlong@ios.ac.cn; zhangxd20@ios.ac.cn

ORCID: 0009-0007-5508-3404; 0000-0002-8979-7014; 0009-0005-0451-7810 *Corresponding author

ABSTRACT

RAG attribution is often scored per document although a response may have several alternative minimal causes. We formalize attribution as enumeration of all inclusion-minimal sufficient document sets and prove a boundary for grand-coalition leave-one-out (LOO): under monotone replay, LOO recovers the unique cause but returns only the intersection of alternatives. We introduce oracled unique/alternative/non-monotone (U/A/N) strata and audit black-box procedures on HotpotQA and Natural Questions with two 7B generators. Mask counts overstate cost when masks render to identical prompts; correct prompt-level caching removes an apparent advantage of a RAG-specific adaptation over enumeration. An enumerate-first prompt-uniform residual audit preserves perfect U/A recovery and raises source-clustered raw-N violation recall from .388 to .952, but a no-collapse stress test falls to .814. We prove uniform sampling minimax for an adversarial static hidden-witness subproblem. A frozen geometry-aware ordering fails its prespecified holdout gate. Thus a finite unviolated trace is diagnostic, not a monotonicity certificate.

Index Terms— retrieval-augmented generation, context attribution, sufficient causes, leave-one-out, evaluation audit

1. INTRODUCTION

RAG combines parametric generation with retrieved evidence [1, 2, 3], but attribution still asks which passages caused a fixed answer. Existing methods assign document scores or citations [4, 5, 6], but a scalar ranking does not identify a *set* of alternative minimal sufficient causes. This distinction matters in poisoning and forensic analysis: removing one path can leave another intact. Minimal sufficient subsets are established in explanation research [7, 8, 9]; our question is narrower: what can exact black-box replay establish for a fixed RAG context?

Contributions. We make four contributions. (1) *Identifiability*: we prove that LOO exactly recovers a unique monotone cause but only the common backbone of alternatives, and that an unviolated finite trace is not a monotonicity certificate. (2) *Benchmarking*: we construct fully enumerated U/A/N strata whose labels come from the realized generator truth table, never the planted template. (3) *Methodology*: we introduce an enumerate-first pipeline that protects all-cause search, deduplicates byte-identical rendered prompts, and spends only the residual physical-call budget on safety auditing; uniform class sampling is minimax in a static hidden-witness

game. (4) *Audit evidence*: corrected accounting removes an apparent method advantage; transfer, no-collapse stress, and two frozen holdout attempts expose where the diagnostic succeeds and where it must abstain.

Relation to prior work. Sufficient-input subsets and rationale benchmarks formalize local evidence [7, 10], while delta debugging minimizes a failure-inducing input [11]. ContextCite fits a KernelSHAP surrogate over context ablations [4, 12]; MIRAGE and SelfCite instead use model internals or citation supervision [5, 6]. Our all-cause endpoint differs from feature ranking and single-subset minimization [13, 14, 15]. Explanation quality also depends on the claim and evaluation target [16, 17]; we therefore make faithfulness assumptions explicit and prefer intervention-backed tests [18, 19, 20, 21, 22, 23, 24]. Post-hoc local surrogates can also be adversarially misleading [25]. Its standard baseline follows the same minimal/maximal duality used by MUS/MCS enumeration [26]. Property testing can reject functions that are far from monotone [27]; our trace proposition is narrower and makes no tolerant- testing lower-bound claim.

Data availability. Code, frozen protocols, benchmark manifests, and replay-derived results are available at <https://github.com/violetxd/CoalitionTrace>.

2. PROBLEM SETUP AND BOUNDARIES

For documents $D = \{1, \dots, n\}$, let $r(S)$ be the complete rendered chat prompt (including the system message and generation marker) obtained by keeping the observed version of documents in S and their contrastive replacement otherwise. Deterministic replay defines $v(S) = \mathbb{I}[\text{NEM}(g(r(S))) = \text{NEM}(y_t)]$, where NEM is a frozen normalization and exact-match rule: a format-only extractor removes an “Answer:” prefix, later lines, and a trailing parenthetical before standard lowercase/article/punctuation normalization. A set C is an inclusion-minimal sufficient cause (relative to the fixed model, decoder, replacement, and target) when $v(C) = 1$ and $v(C') = 0$ for every $C' \subsetneq C$. Let $\mathcal{M}(v)$ contain all such causes. We evaluate exact cause-set match (ACEM), cause-set precision/recall/F1, observed non-monotonicity, and replays.

Proposition 1 (LOO boundary). Write $D_{-i} = D \setminus \{i\}$. If v is monotone and $v(D) = 1$, grand-coalition LOO returns

$$\{i : v(D_{-i}) = 0\} = \bigcap_{C \in \mathcal{M}(v)} C. \quad (1)$$

Hence LOO returns an inclusion-minimal sufficient cause exactly in the unique regime, $|\mathcal{M}(v)| = 1$. With alternatives it returns only

their common backbone. The proof follows because every sufficient set contains a minimal cause: deleting i fails exactly when every minimal cause contains i .

Corollary (one-query ambiguity test). Let B be the LOO set above. Under the same assumptions, $v(B) = 1$ iff $|\mathcal{M}(v)| = 1$. Thus at most $n+2$ replays (full, n deletions, and B) classify U versus A and recover the U cause. If B is sufficient it contains a minimal cause, which must equal the intersection and hence every minimal cause; the converse follows from Proposition 1. This certificate is invalid without monotonicity.

This motivates three oracle strata, all requiring $v(\emptyset) = 0$ and $v(D) = 1$: U has monotone v and one minimal cause; A has monotone v and at least two; N contains a Boolean cover-edge violation.

Proposition 2 (trace non-certifiability). Assume a partial trace over exact rendered-prompt classes has a monotone completion. If an entirely unqueried class represented by x has either a queried insufficient strict superset or a queried sufficient strict subset, then a non-monotone completion produces the identical trace: assign every mask in that whole unqueried class the opposite of the value forced by monotonicity. Thus absence of an observed violation cannot certify monotonicity for such a trace.

Proposition 3 (static hidden-witness minimax). After a fixed prefix, suppose M classes remain, exactly K are immediate witnesses, every non-witness returns one common miss, misses cannot create new witnesses, and all K -subsets are admissible. Any randomized adaptive audit of $b \leq M$ classes has worst-case detection at most $V = 1 - \binom{M-K}{b} / \binom{M}{b}$; uniform sampling without replacement attains V for every witness set. Indeed, under a uniformly hidden K -set, any fixed all-miss query set of size at most b intersects it with probability at most V , so averaging yields one instance no easier than V ; uniform sampling has that probability pointwise. In a real sequential audit, dynamic witnesses can only make the fixed-prefix expression a lower bound, not a global minimax claim.

3. METHODOLOGY

Design objective. CoalitionTrace separates three questions that a single attribution score cannot answer: (i) which minimal causes have been recovered, (ii) whether monotonicity has been refuted, and (iii) what the trace actually certifies. Figure 1 shows the resulting composition. Cause enumeration receives the task budget first; only its unused *physical-call* budget is audited. Thus safety evidence cannot silently replace the all-cause objective it is meant to qualify.

Contract and state. We freeze D , slot-preserving replacements, target y_t , generator g , renderer r , scorer, threshold, and total budget; a cause is meaningful only under this intervention. The persistent state is the exact prompt cache H , verified causes $\widehat{\mathcal{M}}$, and exclusion frontier F . An evidence ledger stores masks, physical calls, multiplicities, values, and the stopping reason.

Phase I: scope validation and cheap ambiguity diagnosis. We replay $\emptyset$ and D first. If $v(\emptyset) = 1$ there is no document-level necessity; if $v(D) = 0$ the target is absent at the observed endpoint. Both are out of scope. We then query all D_{-i} and form the LOO backbone $B = \{i : v(D_{-i}) = 0\}$. If monotonicity is justified, one additional query $v(B)$ implements the corollary’s U/A test: $v(B) = 1$ identifies the unique minimal cause, whereas $v(B) = 0$ exposes alternative paths. Without that assumption, B is only a diagnostic initialization.

Phase II: protected all-cause search. Initialize $F = \{\emptyset\}$. For $E \in F$, inspect $D \setminus E$. If it is insufficient, monotonicity prunes the branch; otherwise deterministic chunk and member deletion yields C . Accept it only after replay checks $v(C) = 1$ and $v(C \setminus \{i\}) = 0$

for all $i \in C$, then branch to $E \cup \{i\}$ for each $i \in C$. This hitting-set invariant leaves a frontier node disjoint from every undiscovered cause. Exhausting F is conditionally complete under monotonicity; budget exhaustion returns verified causes with incomplete status. This is a standard enumerator [8, 9]; our contribution is its protected composition with exact accounting and safety grading.

Phase III: residual audit and evidence grade. Let $B_{\text{residual}} = B_{\text{total}} - B_{\text{enumeration}}$. Only after Phase II stops does the audit sample unseen classes in $2^D / \sim$ uniformly without replacement, propagate values to byte-identical masks, and test observed comparable pairs. The first $L \subset U$ with $v(L) = 1, v(U) = 0$ is a replayable refutation. The grade is *exact-monotone* only after required lattice exhaustion, *conditional-complete* after frontier exhaustion under monotonicity, *violated* with a witness, or *budget-limited*; a clean partial trace abstains.

Algorithmic advantages. The composition provides four concrete guarantees: (1) residual auditing cannot consume cause-search queries at the same total budget; (2) H charges physical calls, not duplicate masks, without semantic merging; (3) violations carry explicit witnesses while clean partial traces never become global claims; and (4) the ledger reconstructs every cause check, cache hit, witness, and cost. The output is evidence plus scope, not only a scalar ranking.

3.1. Exact replay accounting

Define $S \sim T$ iff $r(S) = r(T)$ byte-for-byte. This is the conservative exact request equivalence used by our cache, with model, decoder, scorer, and fixed randomness held constant; token-ID equivalence could merge additional byte-distinct prompts and is not claimed here. **Proposition 4 (quotient accounting).** Any adaptive policy depending only on its logical mask/response transcript can be simulated exactly with one call per distinct class in $\mathcal{Q} = 2^D / \sim$. Under b uniform mask draws, class q is included with probability $1 - \binom{2^n - m_q}{b} / \binom{2^n}{b}$, biased by its multiplicity m_q . For a *static* single witness class, when its identity does not change and every miss yields the same all-miss transcript, no adaptive policy using c distinct class calls guarantees detection above $c/|\mathcal{Q}|$: inclusion marginals sum to c , and class-uniform sampling attains the bound. Cache induction, avoidance counting, and averaging prove the claims. The quotient is lossless for calls, not a cause lattice, so queried values are propagated to all member masks before subset comparisons.

4. EVALUATION

We ask whether LOO exhibits the predicted U/A boundary; whether prompt-level accounting changes method comparisons; whether residual auditing improves N-witness recall without sacrificing cause recovery; and whether development gains survive transfer, no-collapse stress, and frozen holdouts.

4.1. Benchmark and realized ground truth

We start from public PoisonedRAG target questions [28] for HotpotQA and Natural Questions [29, 30]. Each $n = 8$ context contains contrastive document slots that instantiate one path (U), two disjoint or overlapping paths (A), or a suppressor/restorer pattern (N). We exhaust all 2^8 masks with greedy decoding. We retain a source only when its realized table satisfies the exact stratum definition; all failures and the screening funnel are retained in a hashed artifact provided in the submitted supplementary material. A source query, not a template or mask, is the sampling unit. The *primary registry* frozen




Fig. 1. CoalitionTrace separates the frozen causal contract, exact request cache, protected cause enumeration, residual safety audit, and evidence grade. All stages share H and write an auditable ledger. The highlighted properties are design invariants, not empirical side effects.

test split contains 70 sources per dataset; later news, all-active, and holdout audits use separately reported pools. Models are Qwen2.5-7B-Instruct and Mistral-7B-Instruct-v0.3 at pinned revisions.

4.2. Baselines, metrics, and protocol

Methods are LOO; hierarchical deletion; ten-seed repeated greedy deletion; ten-seed random masks; constrained KernelSHAP as a ContextCite-style surrogate; a standard monotone exclusion enumerator; and a RAG adaptation that adds an LOO backbone fast path. The post-hoc residual audit is evaluated separately from the frozen seven-method comparison, with ten seeds. All share the same SHA-256 prompt cache. The primary small- n budget is 48 unique prompts. Exact tables supply truth, not a deployable method. We bootstrap source queries (10,000 resamples); stochastic seeds are averaged within source and are never treated as samples. This intervention-based design tests faithfulness directly rather than treating attention or plausibility as a causal explanation [19, 18]; it also differs from training-data influence analysis [31].

Oracle and construction validity. Every stored table is checked for all 2^n masks, target normalization, $v(\emptyset)$, and $v(D)$. Method-blind selection keeps the first case in a frozen candidate order. An independent implementation re-derived labels and causes for 224 candidate tables and found zero mismatches. Planted templates are not truth: 79–90% of realized A cases contain a cause crossing planted paths, while only 4–21% exactly match the planted cause set. Cells contain two to at least four causes, of size one to five; hence we score realized cause-set equality.

Reproducibility and cost. Model revisions, tokenizer chat templates, case order, replacements, and greedy decoding (at most 16 new tokens) are pinned. Each replay stores mask ID, raw/normalized answer, prompt SHA-256, revision, tokens, and class size. The evaluator verifies sufficiency and member-necessity against the exhaustive table; invalid subsets receive no credit. Stochastic baselines use one fixed 10-seed panel.

4.3. Attribution and accounting results

LOO has a sharp empirical boundary. Across completed registry cells, LOO obtains U ACEM 1.00 and A ACEM 0, exactly matching

Proposition 1. The standard enumerator and RAG adaptation obtain A ACEM 1.00 in the completed fair cells. Hierarchical deletion and KernelSHAP also return at most one subset and therefore have A ACEM 0 under the all-causes endpoint; this is an endpoint mismatch, not evidence that their single returned subset is invalid. On monotone A, every Adapt and StdEnum return is a member of the exhaustive $\mathcal{M}(v)$ (cause precision 1.00; invalid-return rate 0 in all four cells).

Correct caching collapses the claimed advantage. In registry A cases, the complete 256-mask tables contain only 15.7–19.8 distinct rendered prompts on average (6.1–7.7%). Under the original mask-count budget, the Qwen adaptation exceeded repeated deletion by 20.5/20.2 ACEM points on Hotpot/NQ; under unique-prompt accounting the gaps shrink to 9.5/9.6 points, below the preregistered 10-point criterion. Random-mask ACEM changes from .00/.00 to 1.00/.920 because equivalent masks no longer consume fictitious generator calls. With unique-prompt accounting, Mistral A ACEM is 1.00 for the adaptation, .913/.931 for repeated deletion, and .981/.983 for random masks on Hotpot/NQ. The standard enumerator matches 1.00 while using fewer prompts in every cell. These results reject the preregistered method-superiority criterion and show that this A slice is close to saturated. Source-level Adapt–Repeated gains are .095, .096, .087, and .069 in Table 1 order; all are below the prespecified .10 margin, and their source-paired intervals are reported in the released artifact. Adapt–StdEnum ACEM is zero in every cell (Holm-adjusted exact $p = 1$), while StdEnum uses fewer prompts. The frozen news family again gives both enumerators A ACEM 1.00 ($n = 38$ –53). At $n = 10$, only 1.7–2.1% of 1,024 masks are distinct prompts and both remain at 1.00. The all-active control has $|\mathcal{Q}| = 256$ but yields only 6/5/3/4 A sources; it is a no-collapse stress test, not an accuracy benchmark.

4.4. Safety-audit and generalization results

Residual auditing is diagnostic, not confirmatory. Standard-Enum already observes violations in .267–.448 of raw N cases. Prompt-uniform residual Audit raises ten-seed mean recall to .922–.980. Averaging model outcomes within each dataset/source cluster gives .952 over 130 source clusters (222 model-cases; bootstrap 95% CI [.930, .972]) while retaining U/A ACEM 1.00. Witness

Table 1. A-stratum ACEM / mean unique prompts at budget 48 (n_A in parentheses). Repeated, Random, and Audit average ten seeds; StdEnum is the task-matched enumerator and Adapt adds its LOO fast path.

Data	Model	LOO	Repeated	Random	StdEnum	Adapt	Audit
Hotpot (58)	Qwen	.00 (5.9)	.905 (13.0)	1.00 (15.7)	1.00 (12.4)	1.00 (13.0)	1.00 (15.7)
NQ (49)	Qwen	.00 (6.1)	.904 (14.2)	.920 (18.4)	1.00 (14.2)	1.00 (14.6)	1.00 (18.4)
Hotpot (54)	Mistral	.00 (5.9)	.913 (13.1)	.981 (15.4)	1.00 (11.9)	1.00 (12.9)	1.00 (15.4)
NQ (52)	Mistral	.00 (5.9)	.931 (13.4)	.983 (15.8)	1.00 (12.8)	1.00 (14.1)	1.00 (15.8)




Fig. 2. Accounting audit. (a) Unique-prompt charging reduces the apparent adaptation gap. (b) Registry prompts collapse strongly; the all-active control does not. Q/M denote Qwen/Mistral.




Fig. 3. Residual-audit budget frontier. Lines are cell means and bands are minima-maxima; enumeration is never preempted.

soundness implies no false flags, confirmed on all 429 monotone U/A model-cases. Its paired gain over the identical enumeration prefix is .563 [.495,.630]. On detected N cases, the witness appears after 4.05 residual prompt queries on average. Figure 3 shows the budget tradeoff: at 12,16,24,32,48 prompts, mean A ACEM is .545,.967,.991,.995,1.00 and mean N recall is .421,.635,.734,.841,.949.

The wrapper was designed after inspecting the frozen comparison, so these are explicitly post-hoc results. We freeze it for a post-hoc transfer check on the separate news-wording family: N recall is .950 [.929,.968] over 123 clusters (207 model-cases), and A ACEM remains 1.00. The all-active control, where $|Q| = 256$, falls to .814 [.753,.869] over 75 clusters (142 model-cases); this negative stress result prevents a near-certification interpretation. Ablations give .797/.667/.935 for deterministic upward/downward/cover audits and .948 for randomized cover. Thus the robust benefit is protected enumeration plus residual budget, not a uniquely superior boundary heuristic.

Across 1,000 offline permutations per N case, fixed-prefix and sequential recall are .920872 and .920959 (571 model-cases;

138 clusters), with difference .000087 [−.000030,.000495] and no dynamic-only detection.

We froze nearest ordering before a two-dataset/two-model hold-out. Method-blind selection retained 67/64 N cases in MSMARCO but only 15/16 in NQ, below the 30-per-cell floor. Descriptively, recall was .835 [.781,.886], with gains over uniform of .038 [.004,.072] (below the .05 margin) and shuffled cover of .028 [−.009,.064]; one cell was negative. A separately preregistered feasibility pilot on source-disjoint SQuAD and WebQuestions data [32, 33] also stopped before method evaluation: N yields 14/13/19/14 all missed its 20-source floor; no final replay or retuning followed. These blind failures accord with Proposition 3: geometry may help under a validated distribution but cannot improve the adversarial static guarantee. All methods make zero unconditional completeness claims on N; by Proposition 2, an unviolated trace is inconclusive.

5. LIMITATIONS AND CONCLUSION

The deployment rule is assumption-sensitive: validate D and $\emptyset$; under defensible monotonicity, use $v(B)$ to distinguish a unique cause from alternatives, then enumerate. Otherwise only a reversal is decisive and a clean partial trace abstains. Uniform class auditing is a small- n method, and cost must separate masks, unique prompts, and tokens. Results are local to constructed contexts, fixed replacements, exact match, and two 7B models; they do not estimate natural U/A/N prevalence. Nevertheless, they support a defensible practice: protect all-cause search, audit only with residual budget, and never promote an unviolated trace to a certificate.

6. ACKNOWLEDGMENT

OpenAI Codex assisted code implementation, figure production, and language editing. The authors designed the study, ran and verified the experiments, and take responsibility for the manuscript and its claims.

7. REFERENCES

- [1] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, 2020, vol. 33, pp. 9459–9474.
- [2] Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih, “Dense passage retrieval for open-domain question answering,” in *Proc. EMNLP*, 2020, pp. 6769–6781.
- [3] Gautier Izacard and Edouard Grave, “Leveraging passage retrieval with generative models for open domain question answering,” in *Proc. EACL*, 2021, pp. 874–880.
- [4] Benjamin Cohen-Wang, Harshay Shah, Kristian Georgiev, and Aleksander Madry, “ContextCite: Attributing model generation to context,” in *Advances in Neural Information Processing Systems*, 2024, vol. 37.
- [5] Jirui Qi, Gabriele Sarti, Raquel Fernández, and Arianna Bisazza, “Model internals-based answer attribution for trustworthy retrieval-augmented generation,” in *Proc. EMNLP*, 2024, pp. 6037–6053.
- [6] Yung-Sung Chuang, Benjamin Cohen-Wang, Zejiang Shen, Zhaofeng Wu, Hu Xu, Xi Victoria Lin, James R. Glass, Shang-Wen Li, and Wen-Tau Yih, “SelfCite: Self-supervised alignment for context attribution in large language models,” in *Proc. ICML*, 2025, pp. 10839–10858.
- [7] Brandon Carter, Jonas Mueller, Siddhartha Jain, and David Gifford, “What made you do this? understanding black-box decisions with sufficient input subsets,” in *Proc. AISTATS*, 2019, pp. 567–576.
- [8] Alexey Ignatiev, Nina Narodytska, and Joao Marques-Silva, “Abduction-based explanations for machine learning models,” in *Proc. AAAI*, 2019, vol. 33, pp. 1511–1519.
- [9] Joao Marques-Silva, Thomas Gerspacher, Martin C. Cooper, Alexey Ignatiev, and Nina Narodytska, “Explanations for monotonic classifiers,” in *Proc. ICML*, 2021, pp. 7469–7479.
- [10] Jay DeYoung, Sarthak Jain, Nazneen Fatema Rajani, Eric Lehman, Caiming Xiong, Richard Socher, and Byron C. Wallace, “ERASER: A benchmark to evaluate rationalized NLP models,” in *Proc. ACL*, 2020, pp. 4443–4458.
- [11] Andreas Zeller and Ralf Hildebrandt, “Simplifying and isolating failure-inducing input,” *IEEE Transactions on Software Engineering*, vol. 28, no. 2, pp. 183–200, 2002.
- [12] Scott M. Lundberg and Su-In Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, 2017, vol. 30.
- [13] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin, “Why should I trust you?: Explaining the predictions of any classifier,” in *Proc. KDD*, 2016, pp. 1135–1144.
- [14] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin, “Anchors: High-precision model-agnostic explanations,” in *Proc. AAAI*, 2018, vol. 32, pp. 1527–1535.
- [15] Ruth C. Fong and Andrea Vedaldi, “Interpretable explanations of black boxes by meaningful perturbation,” in *Proc. ICCV*, 2017, pp. 3429–3437.
- [16] Tim Miller, “Explanation in artificial intelligence: Insights from the social sciences,” *Artificial Intelligence*, vol. 267, pp. 1–38, 2019.
- [17] Zachary C. Lipton, “The mythos of model interpretability,” *Communications of the ACM*, vol. 61, no. 10, pp. 36–43, 2018.
- [18] Alon Jacovi and Yoav Goldberg, “Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness?,” in *Proc. ACL*, 2020, pp. 4198–4205.
- [19] Sarthak Jain and Byron C. Wallace, “Attention is not explanation,” in *Proc. NAACL-HLT*, 2019, pp. 3543–3556.
- [20] Sarah Wiegrefe and Yuval Pinter, “Attention is not not explanation,” in *Proc. EMNLP-IJCNLP*, 2019, pp. 11–20.
- [21] Sofia Serrano and Noah A. Smith, “Is attention interpretable?,” in *Proc. ACL*, 2019, pp. 2931–2951.
- [22] Chih-Kuan Yeh, Cheng-Yu Hsieh, Arun Suggala, David I. Inouye, and Pradeep K. Ravikumar, “On the (in)fidelity and sensitivity of explanations,” in *Advances in Neural Information Processing Systems*, 2019, vol. 32.
- [23] Amirata Ghorbani, Abubakar Abid, and James Zou, “Interpretation of neural networks is fragile,” in *Proc. AAAI*, 2019, vol. 33, pp. 3681–3688.
- [24] Sara Hooker, Dumitru Erhan, Pieter-Jan Kindermans, and Been Kim, “A benchmark for interpretability methods in deep neural networks,” in *Advances in Neural Information Processing Systems*, 2019, vol. 32.
- [25] Dylan Slack, Sophie Hilgard, Emily Jia, Sameer Singh, and Himabindu Lakkaraju, “Fooling LIME and SHAP: Adversarial attacks on post hoc explanation methods,” in *Proc. AIES*, 2020, pp. 180–186.
- [26] Mark H. Liffiton, Alessandro Previti, Ammar Malik, and Joao Marques-Silva, “Fast, flexible MUS enumeration,” *Constraints*, vol. 21, pp. 223–250, 2016.
- [27] Oded Goldreich, Shafi Goldwasser, Eric Lehman, Dana Ron, and Alex Samorodnitsky, “Testing monotonicity,” *Combinatorica*, vol. 20, no. 3, pp. 301–337, 2000.
- [28] Wei Zou, Runpeng Geng, Binghui Wang, and Jinyuan Jia, “PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models,” in *Proc. USENIX Security*, 2025, pp. 3827–3844.
- [29] Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning, “HotpotQA: A dataset for diverse, explainable multi-hop question answering,” in *Proc. EMNLP*, 2018, pp. 2369–2380.
- [30] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, et al., “Natural questions: A benchmark for question answering research,” *Transactions of the Association for Computational Linguistics*, vol. 7, pp. 453–466, 2019.
- [31] Pang Wei Koh and Percy Liang, “Understanding black-box predictions via influence functions,” in *Proc. ICML*, 2017, vol. 70, pp. 1885–1894.
- [32] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang, “SQuAD: 100,000+ questions for machine comprehension of text,” in *Proc. EMNLP*, 2016, pp. 2383–2392.
- [33] Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang, “Semantic parsing on Freebase from question-answer pairs,” in *Proc. EMNLP*, 2013, pp. 1533–1544.